

ondia

## Kubernetes Networking



ondia


## AGENDA

- ▶ Cluster Networking
- ▶ Services
- ▶ Service Types
- ▶ Labels and loose coupling

1

## Cluster Networking

## Cluster Networking




## Cluster Networking

There are 4 distinct networking problems to address:

1. container-to-container communications:
2. Pod-to-Pod communications:
3. Pod-to-Service communications:
4. External-to-Service communications:


2

## Services




## Services

Pods are not reliable



## Services

A **Service** offers a single **DNS entry** for a containerized application managed by the Kubernetes cluster




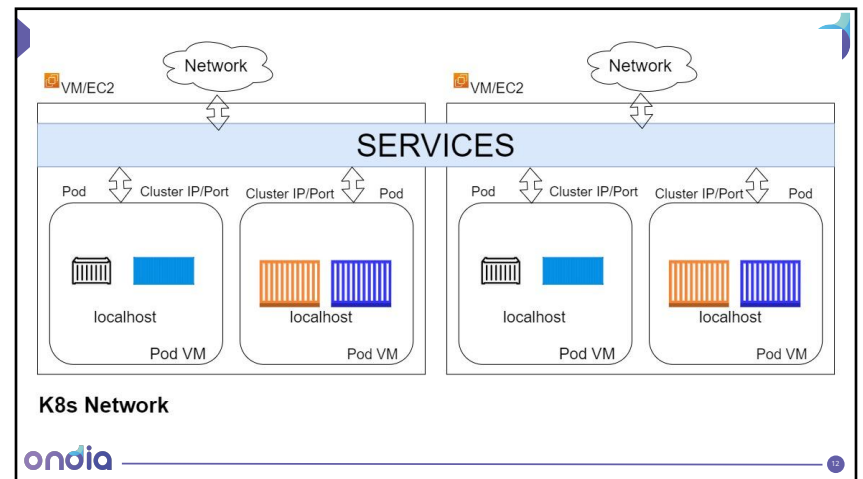

## Services

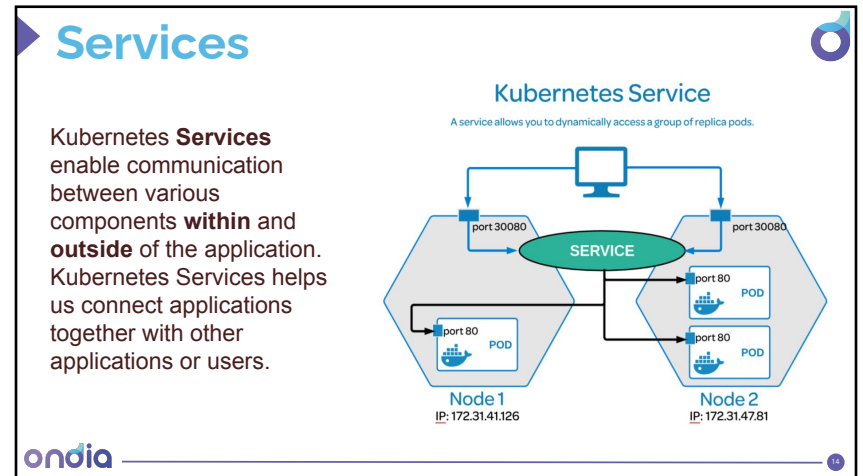
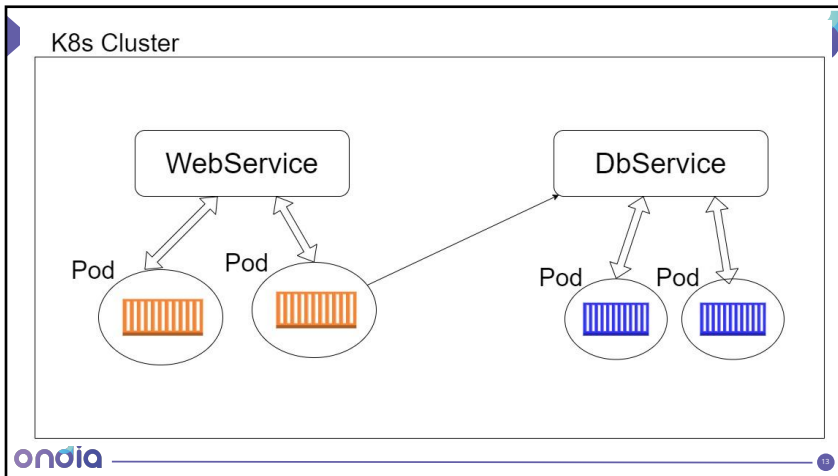
The **Service** is associated with the Pods, and provides them with a stable IP, DNS and port. It also **loadbalances** requests across the Pods.

**Service** logically groups Pods and defines a policy to access them.

This grouping is achieved via **Labels** and **Selectors**.



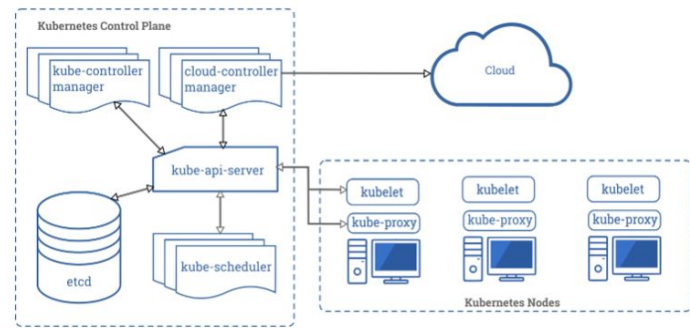




## kube-proxy

- Each cluster node runs a daemon called **kube-proxy**
- **kube-proxy** is responsible for **implementing the Service configuration** on behalf of an administrator or developer
- For each new Service, on each node, **kube-proxy** configures **iptables** rules to capture the traffic for its **ClusterIP** and forwards it to one of the Service's endpoints.
- When the Service is removed, **kube-proxy** removes the corresponding **iptables** rules on all nodes as well.


## Control Plane Components





## Service Discovery


- Kubernetes has an add-on for **DNS**, which creates a DNS record for each Service and its format is




- Services within the same Namespace find other Services just by their names.
- If we add a Service **redis-master** in **my-ns** Namespace, all Pods in the same **my-ns** Namespace lookup the Service just by its name, **redis-master**.

FQDN: fully qualified domain name

**my-ns**



**test-ns**



To connect to the "Game pod" and "db pod":

From "web pod" -> "Game pod" -> `hostname: game-service.my-ns:port`  
`game-service:port`

From "web pod" -> "db pod" -> `hostname: db-service.test-ns.svc.cluster.local:port`



3


## Service Types



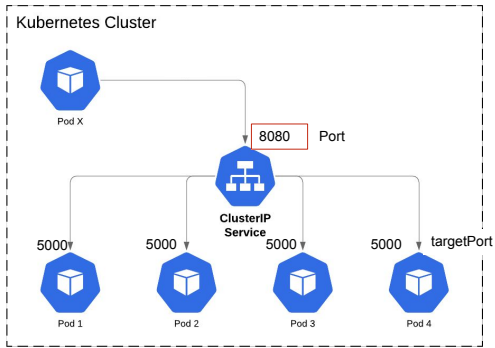
## Service Types

There are 4 major service types:

- ClusterIP (default)
- NodePort
- LoadBalancer
- ExternalName



## Service Types




**ClusterIP:**  
Expose traffic internally


### Example Usecase:

Good for service of  
database & back-end  
apps.

```
apiVersion: v1
kind: Service
metadata:
  name: back-end-svc
  labels:
    app: back-end
spec:
  type: ClusterIP (default)
  selector:
    app: back-end
  ports:
    - port: 3000
      protocol: TCP
      targetPort: 5000
```



## Service Types




**NodePort:**  
Exposes traffic to the outside.

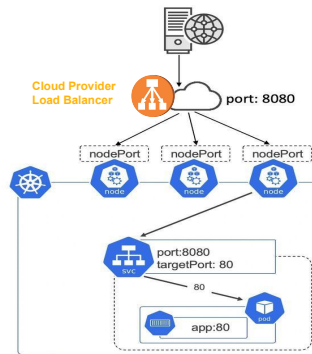
### Example Use Case:

when we want to make our Services which has our app or website accessible from the external world.

```
apiVersion: v1
kind: Service
metadata:
  name: web-flask-svc
  labels:
    app: web-flask
spec:
  type: NodePort
  selector:
    app: web-flask
  ports:
    - nodePort: 30480
      port: 3000
      protocol: TCP
      targetPort: 5000
```



## Service Types




### LoadBalancer:

Exposes traffic outside with load balancing feature.

### Example Use Case:

when we want to load balance our Services which has our app or website accessible from the external world.

```
apiVersion: v1
kind: Service
metadata:
  name: web-flask-svc
  labels:
    app: web-flask
spec:
  type: LoadBalancer
  selector:
    app: web-flask
  ports:
    - port: 3000
      protocol: TCP
      targetPort: 5000
```




## Service Types

### LoadBalancer:

- The **LoadBalancer** *ServiceType* will only work if the underlying infrastructure supports the automatic creation of Load Balancers and have the respective support in Kubernetes, as is the case with the Google Cloud Platform, Azure or AWS.
- If no such feature is configured, the **LoadBalancer IP** address field is **not populated**, it remains in **Pending** state, but the **Service will still work as a typical NodePort type Service**.

## Service Types



## Service Types

### ExternalName:

Maps the Service to the contents of the ExternalName field (e.g. example.com), by returning a CNAME record with its value.

### Example Use Cases:

to make externally configured services like;

`remote.server.url.com`

available to applications inside the cluster.

## Service Types

```
apiVersion: v1
kind: Service
metadata:
  labels: io.kompose.service: mysql-server
  name: mysql-server
spec:
  type: ExternalName
  externalName: mydb.cbanmzptkrzf.us-east-1.rds.amazonaws.com
```

4

## Labels and loose coupling



## Labels and loose coupling

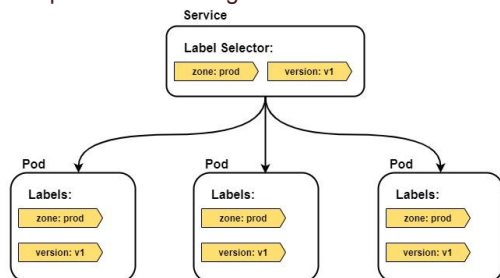


- Labels and Selectors use a **key/value** pair format.
- Pods and Services are loosely coupled via labels and label selectors.
- For a Service to match a set of Pods, and therefore provide stable networking and load-balance, it only needs to match some of the Pods labels.
- However, for a Pod to match a Service, the Pod must match all of the values in the Service's label selector.




## Labels and loose coupling

The figure below shows an example where 3 Pods are labeled as **zone=prod** and **version=v1**, and the Service has a label selector that matches. This Service provides stable networking to all three Pods. It also provides simple load-balancing.

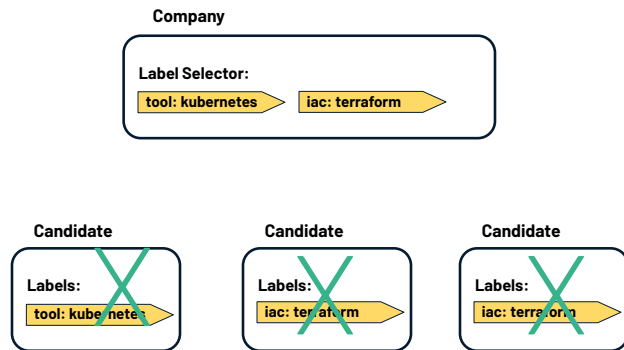


## Labels and loose coupling

The figure below shows an example where the Service does not match any of the Pods. This is because the Service is selecting on two labels, but the Pods only have one of them. The logic behind this is a Boolean AND operation.

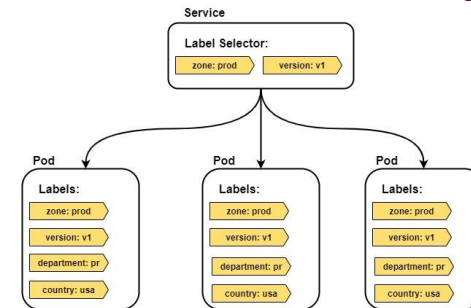


## Labels and loose coupling

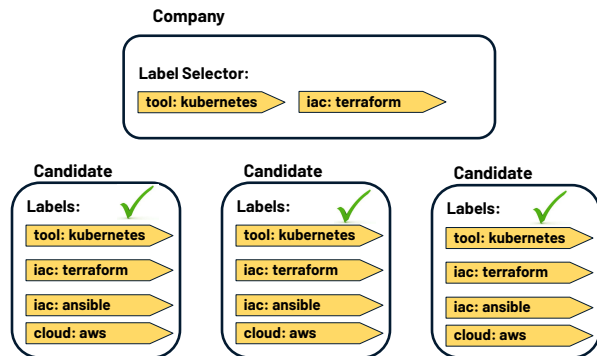


## Labels and loose coupling

This figure shows an example that does work. It doesn't matter that the Pods have additional labels that the Service is not selecting on.



## Labels and loose coupling



**THANKS!**  
Any questions?




5

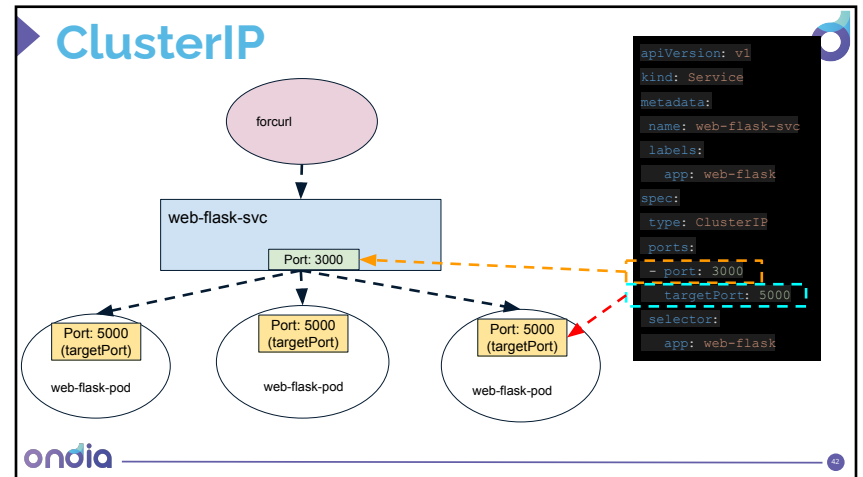
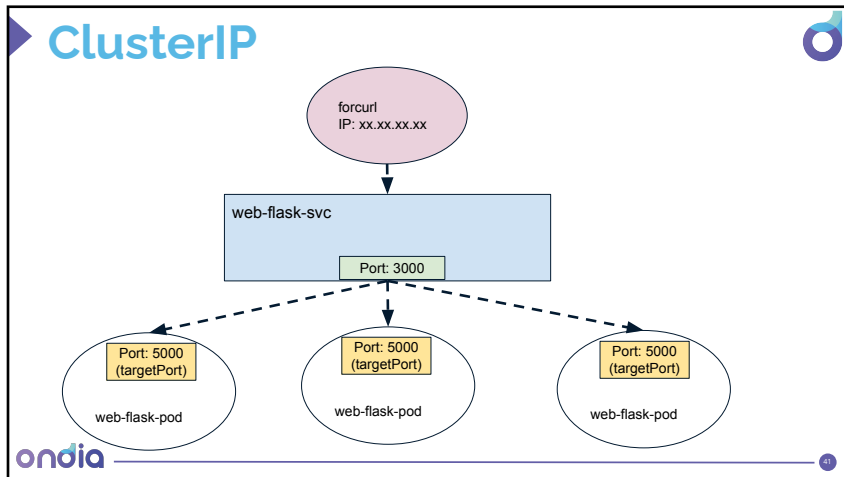
## Kubernetes hands-on-03



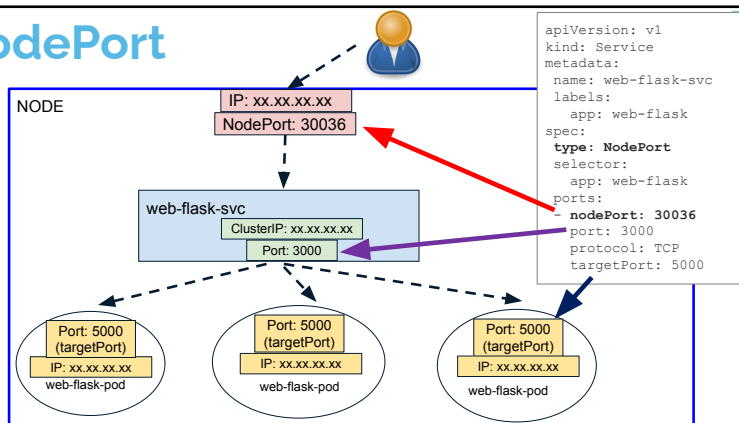
## Pod to Pod Connection

**forcurl** pod can connect to **web** pod directly, if only it knows the IP and port number of the **web** pod.





## NodePort



## Endpoint

